
Predicting Airbnb Listing Prices in Major US Cities

Group 15

Mathieu

Serghei

Samy

Exchange students from ESIEA

Professors : Ajan Aye Hninn Khine and Ajan Naveed Sultan

Project type: Regression
Target variable: log_price
Dataset: AirBnB Listings in Major US Cities (Deloitte ML, Kaggle)
GitHub: <https://github.com/Lolilop668/Final-Project.git>

Submitted July 7, 2026

Contents

1	Project title and members	3
2	Problem definition and goals	3
2.1	Problem statement	3
2.2	Target variable	3
2.3	Why the problem matters	3
2.4	Expected outcome	3
3	Dataset description	4
3.1	Two files: train and test	4
3.2	Cities covered	4
3.3	Features	4
3.4	Limitations of the data	4
4	Preprocessing and cleaning	5
4.1	Train and validation split	5
4.2	Dropping unused columns	5
4.3	Missing values	5
4.4	Encoding and scaling	5
5	Data visualization and exploratory data analysis	5
5.1	Distribution of the numeric features	5
5.2	Where the listings are	6
5.3	Distribution of the categorical features	7
5.4	The target variable	8
5.5	Price by city and by room type	9
5.6	Correlation with the target	9
6	Feature engineering	10
6.1	Feature selection	10
6.2	Encoding and scaling as engineered features	10
6.3	Latitude and longitude as location proxies	11
6.4	What we chose not to do	11
7	Machine learning model training	11
7.1	Models	11
7.2	Training setup	11

7.3	Why these models and how we chose	11
8	Evaluation metrics and results	12
8.1	Metrics	12
8.2	Validation results	12
8.3	Why we kept the Random Forest	12
8.4	Validation diagnostics	12
8.5	Feature importance	13
8.6	Predictions on the test file	14
8.7	Honest reading of the results	15
9	Responsible AI practices	15
9.1	Privacy and personal data	15
9.2	Bias and fairness	15
9.3	Data leakage	15
9.4	Honest reporting	15
9.5	Real-world risks	15
9.6	Use of AI tools	16
10	Conclusion and future work	16
10.1	Main findings	16
10.2	What we learned	16
10.3	Future work	16
11	Team contribution	16
	Acknowledgements	17
	References	17

1 Project title and members

Project title: Predicting Airbnb Listing Prices in Major US Cities.

Group: 15.

Members: Mathieu, Serghei, and Samy. We are three exchange students from ESIEA taking CPE393 at KMUTT during the summer 2026 semester.

GitHub repository: <https://github.com/Lolilop668/Final-Project.git>

Dataset: AirBnB Listings in Major US Cities, published for a Deloitte Machine Learning competition and hosted on Kaggle: <https://www.kaggle.com/datasets/rudymizrahi/airbnb-listings-in-major-us>

Every figure in this report comes from our notebook, and the numbers we quote match the outputs of the code we submitted.

2 Problem definition and goals

Setting the right nightly price on Airbnb is not easy. Hosts want a number that fills their calendar without leaving money on the table, and guests want to know whether a listing is a fair deal. Airbnb offers a price suggestion tool, but the logic behind it is hidden. We wanted to see how far we could get with public features alone and a few standard machine learning models.

2.1 Problem statement

Given the characteristics of an Airbnb listing, such as how many guests it fits, how many bedrooms and bathrooms it has, what type of room it is, and where it sits, can we predict its price? This is a supervised regression task. The value we predict is a continuous number, not a category.

2.2 Target variable

The target is `log_price`, the natural logarithm of the nightly price in US dollars. Raw prices are heavily right-skewed, with a small number of very expensive listings pulling the average up. Taking the logarithm compresses that long tail and gives a distribution that is much closer to symmetric, which suits the regression models better.

2.3 Why the problem matters

A working price model has a few practical uses. A new host with no pricing history could get a starting estimate. A guest could check whether a listing looks overpriced for its size and area. An analyst could study which features move the price the most across different cities. Our goal is not to build a production pricing engine, but to run a clean end-to-end project and to be honest about what the model can and cannot do.

2.4 Expected outcome

We expected location and capacity to be the strongest signals, and we expected the tree-based models to beat the linear ones because price does not scale in a simple straight line with the features. The results confirmed both expectations, which we discuss in Sections 8 and 9.

3 Dataset description

3.1 Two files: train and test

The Kaggle dataset comes as two CSV files. The training file, `train.csv`, has 74,111 listings and 29 columns, and it includes the target `log_price`. The test file, `test.csv`, has 25,458 listings and 28 columns, and it does **not** include `log_price`. This shapes our whole workflow: all analysis, training, and model comparison happens on `train.csv`, and `test.csv` is used only at the very end to generate predictions for unseen listings.

3.2 Cities covered

The listings come from six US cities. The split is uneven, which matters later when we read city-level effects:

Table 1: Number of listings per city (training file).

City	Listings
New York City (NYC)	32,349
Los Angeles (LA)	22,453
San Francisco (SF)	6,434
Washington DC (DC)	5,688
Chicago	3,719
Boston	3,468

New York and Los Angeles together make up about three quarters of the rows. Boston and Chicago are the smallest groups.

3.3 Features

The columns fall into a few groups. There are numeric features such as `accommodates`, `bathrooms`, `bedrooms`, `beds`, `number_of_reviews`, and `review_scores_rating`. There are categorical features such as `property_type`, `room_type`, `bed_type`, `cancellation_policy`, and `city`. A few columns are booleans stored as text, like `cleaning_fee`, `instant_bookable`, `host_has_profile_pic`, and `host_identity_verified`. The rest are free text or identifiers we did not use for modeling, including `description`, `name`, `amenities`, `thumbnail_url`, the review dates, and `host_response_rate`. Two columns, `latitude` and `longitude`, give the exact location of each listing.

3.4 Limitations of the data

Three limitations stand out. First, the data is a snapshot from around 2017 and 2018, so the prices are out of date and should not be read as current market values. Second, the geography is narrow, only six large US cities, so nothing here transfers to smaller towns or to other countries. Third, several columns have missing values, which we handle in the next section. The biggest gaps are in `host_response_rate` (about 25% missing) and `review_scores_rating` (about 23% missing, because new listings have no reviews yet).

4 Preprocessing and cleaning

We wrote the cleaning steps as a single function so that the training data and the test data go through exactly the same transformations. This is the part of the pipeline where data leakage usually creeps in, so we fit everything (the encoders and the scaler) on the training data only.

4.1 Train and validation split

Because `test.csv` has no target, we cannot score the model on it. So we split `train.csv` into an 80% training part and a 20% validation part, using a fixed `random_state` of 42 so the split is always the same. The validation part is what we use to compare models and to check that the final model behaves well.

4.2 Dropping unused columns

We removed columns that a simple model cannot use well: identifiers (`id`), free text (`description`, `name`, `amenities`), the image link (`thumbnail_url`), and the date columns (`first_review`, `last_review`, `host_since`). We also dropped `zipcode`, `neighbourhood`, and `host_response_rate`. The first two overlap with the latitude and longitude we already keep, and `host_response_rate` is missing for a quarter of the rows and stored as a string with a percent sign.

4.3 Missing values

For the numeric features we keep (`bathrooms`, `bedrooms`, `beds`, `review_scores_rating`), we fill the gaps with the median. The median is safer than the mean here because these columns are skewed. For categorical columns we replace any missing value with the string `missing` so that it becomes its own category instead of breaking the encoder.

4.4 Encoding and scaling

The categorical columns are text, and scikit-learn models need numbers, so we used `LabelEncoder` to turn each category into an integer. The encoders are fitted on the training data. When we later apply them to the test file, any category the encoder has never seen is mapped to a safe default so the pipeline does not crash. We then scaled the numeric features with `StandardScaler`, again fitting the scaler on the training split and reusing it on the validation and test data. Scaling helps the linear models and the Ridge penalty behave sensibly; the tree models do not need it, but applying the same scaler to all four keeps the code simple.

5 Data visualization and exploratory data analysis

Before modeling we spent time looking at the training data. The goal was to understand the shape of each variable, spot problems, and get a first idea of which features carry price information.

5.1 Distribution of the numeric features

Figure 1 shows histograms of the six main numeric columns. Most listings are small. Capacity is usually 2 to 4 guests, and one or two bedrooms is the norm. The `number_of_reviews` column has

a very long tail, with a handful of listings collecting hundreds of reviews while most have only a few. The rating column is squeezed up near 90 to 100, so on its own it separates listings poorly.




Figure 1: Distribution of the numeric variables in the training set. Capacity, bedrooms, bathrooms, and beds are right-skewed; ratings cluster near the top of the scale.

5.2 Where the listings are

To get a feel for the geography we plotted the listings on a map. In the notebook this is an interactive Folium map with clustered markers, which is the best way to zoom into a city and inspect single listings. For the report we show a static version in Figure 2, with each point colored by its log price. The six cities appear as separate clusters, and within each city the more expensive listings tend to group in the central and coastal areas.




Figure 2: Listing locations colored by `log_price` (sample of 4,000 points). Each cluster is one city; warmer colors are pricier listings. The notebook contains the full interactive Folium version.

5.3 Distribution of the categorical features

Figures 3, 4, and 5 show the three categorical features that matter most. Apartments dominate the property types, houses are a distant second, and everything else is rare. For room type, entire homes and private rooms make up almost the whole dataset, while shared rooms are uncommon. The city plot repeats what the table in Section 3 already showed: NYC and LA carry most of the weight.




Figure 3: Distribution of property type. Apartments and houses account for the large majority of listings.




Figure 4: Distribution of room type. Shared rooms are rare, which is worth remembering when we read the price boxplots.




Figure 5: Number of listings per city. The class imbalance is strong, with NYC and LA far ahead of the rest.

5.4 The target variable

Figure 6 shows the distribution of `log_price`. After the log transform the shape is close to a bell curve, centered a little below 5 in log units. This is exactly why the log price is easier to model than the raw dollar price. There is a small spike at the very bottom of the range that corresponds to a few suspicious near-zero prices.




Figure 6: Distribution of the target variable `log_price`. The log transform gives a roughly symmetric distribution suitable for regression.

5.5 Price by city and by room type

Figure 7 splits the price by city and by room type. San Francisco and Boston have the highest median prices, and Chicago the lowest. The room type effect is clear and intuitive: entire homes cost more than private rooms, which cost more than shared rooms. These two plots already suggest that room type and location will be strong predictors.




Figure 7: `log_price` by city (left) and by room type (right). Location and room type both shift the price in a visible and sensible way.

5.6 Correlation with the target

Figure 8 is the correlation matrix after encoding. Reading the `log_price` row, the strongest linear links are with `room_type` (about -0.61 , because entire homes get a low encoded value and shared rooms a high one) and `accommodates` (about $+0.57$). Then come `bedrooms` ($+0.47$), `beds` ($+0.44$), and `bathrooms` ($+0.36$). Latitude, longitude, and the number of reviews show almost no linear correlation on their own. That does not make them useless: a tree model can still pick up location through non-linear splits, which is exactly what happens in Section 8.




Figure 8: Correlation matrix of the encoded features and the target. Capacity-related features correlate positively with price; room type correlates strongly in the opposite direction because of its encoding.

6 Feature engineering

Our feature engineering is deliberate rather than heavy. For an honest baseline we preferred a small, well-understood feature set over a large one full of noise.

6.1 Feature selection

From the original columns we kept 17 as model inputs: 8 numeric features and 9 categorical features. We dropped the text, the identifiers, and the date columns for the reasons given in Section 4. Keeping the feature set small makes the models faster to train and much easier to interpret.

6.2 Encoding and scaling as engineered features

The main transformations are the label encoding of the categorical columns and the standard scaling of the numeric ones. Both are fitted on the training data only. Encoding turns text categories into integers the models can read. Scaling puts the numeric features on a common range so the linear models and the Ridge penalty treat them fairly.

6.3 Latitude and longitude as location proxies

We chose to keep the raw coordinates instead of the neighbourhood name. Coordinates are numeric, complete, and free of the spelling inconsistencies that plague text location fields. A single neighbourhood column would have needed one-hot encoding into hundreds of sparse columns. The correlation of the coordinates with price is weak on a straight line, but the Random Forest uses them heavily, which tells us they encode location value in a non-linear way.

6.4 What we chose not to do

We did not mine the `amenities` text or the `description` for extra signals. Both are rich but would need natural language processing to use properly, which is beyond the scope of this project. We list this as future work in Section 10.

7 Machine learning model training

The project asks for at least three suitable models. We trained four regression models and compared them on the same data with the same metrics.

7.1 Models

We used Linear Regression as a plain baseline, Ridge Regression as a regularized linear model, a Random Forest Regressor as a non-linear tree-based model, and a Gradient Boosting Regressor as a second tree-based model that builds trees one after another. Linear Regression fits a straight-line relationship between the features and the log price. Ridge does the same but adds a penalty on large coefficients to reduce overfitting. Random Forest builds many independent trees on random subsets of the data and averages them. Gradient Boosting builds trees in sequence, where each new tree tries to fix the errors of the ones before it.

7.2 Training setup

We fit each model on the scaled training split and measured it on the validation split. The Random Forest uses 200 trees with a fixed random seed; the other models use their default settings, which gives a fair and reproducible first comparison. The test file stays untouched until the prediction step, since it has no target to score against.

7.3 Why these models and how we chose

The two linear models give a clear, interpretable reference point and train almost instantly. The two tree models are the natural next step when a straight line is too simple. We did not pick the final model automatically from a single score. We compared all four on the validation set, then chose manually, weighing both performance and interpretability. This keeps the reasoning transparent and avoids selecting a model just because it won one split by a small margin.

8 Evaluation metrics and results

8.1 Metrics

For a regression task the right metrics are the mean absolute error (MAE), the root mean squared error (RMSE), and the coefficient of determination (R^2). MAE is the average size of the error in log units. RMSE punishes large errors more heavily. R^2 says how much of the variance in the price the model explains, where 1.0 is perfect and 0.0 is no better than always guessing the mean.

8.2 Validation results

Table 2 shows the scores of the four models on the validation set. The two linear models land in exactly the same place: with standardized features the default Ridge penalty barely changes the coefficients, so regularization is not what limits them here. Both tree models do clearly better, and the Random Forest edges out Gradient Boosting on all three metrics.

Table 2: Model comparison on the validation set. Lower MAE and RMSE are better; higher R^2 is better.

Model	MAE	RMSE	R^2
Random Forest	0.316	0.433	0.641
Gradient Boosting	0.316	0.413	0.628
Linear Regression	0.369	0.48	0.543
Ridge Regression	0.369	0.48	0.543




Figure 9: MAE, RMSE, and R^2 for the four models on the validation set. The Random Forest is best on all three, with Gradient Boosting close behind.

8.3 Why we kept the Random Forest

Random Forest and Gradient Boosting are close, but we kept the Random Forest as the final model. It had the best validation scores, it trains quickly, and it reports feature importances that help us explain what it learned. Gradient Boosting could likely match or beat it with more tuning, which we note as future work.

8.4 Validation diagnostics

Because the test file has no labels, our real evaluation lives on the validation split, where the true prices are known. Figure 10 shows four diagnostic views of the Random Forest. The top-left scatter of predictions against real values sits close to the diagonal, with more scatter at the two

extremes where data is thin. The top-right plot compares the distribution of predicted prices with the real ones; the predictions are slightly narrower, which is normal for a model that minimizes squared error. The bottom two panels look at the residuals, the difference between real and predicted. They are centered on zero (the mean residual is about -0.002) and show no strong pattern against the predicted value, which is what a well-behaved model should look like.




Figure 10: Random Forest validation diagnostics: predictions vs real values, real vs predicted distribution, the residual distribution, and residuals vs predicted values. The residuals are centered on zero with no obvious trend.

8.5 Feature importance

Figure 11 shows what the Random Forest relied on. Room type is by far the most important feature, at about 37% of the total importance. Longitude and latitude come next, together around 28%, which confirms that location carries a lot of price information even though its linear correlation was weak. Bathrooms follow at about 12%. This ranking matches the story from the EDA: what kind of space you rent and where it is drive most of the price.




Figure 11: Feature importance from the Random Forest. Room type dominates, followed by the geographic coordinates and the number of bathrooms.

8.6 Predictions on the test file

Finally, we applied the Random Forest to the 25,458 listings in `test.csv`, reusing the encoders and the scaler fitted on the training data, and saved the predicted `log_price` for each one to `airbnb_predictions.csv`. We cannot compute MAE, RMSE, or R^2 here, because the file has no true prices. What we can check is whether the predictions look reasonable. Figure 12 shows that the predicted prices on the test file follow almost the same distribution as the predictions on the validation set, which is a good sign that the model is behaving consistently on genuinely unseen data.




Figure 12: Left: distribution of the predicted prices on `test.csv`. Right: the test predictions overlaid on the validation predictions. The two distributions line up closely.

8.7 Honest reading of the results

An MAE of 0.285 in log units means a typical prediction is off by roughly 30% of the true nightly price once converted back to dollars. That is good enough to give a host a ballpark starting figure, but not precise enough to set a real price automatically. An R^2 near 0.70 means the model explains about two thirds of the variation in price, and leaves one third unexplained. A large part of that missing third is demand, seasonality, and listing quality that our features simply do not capture.

9 Responsible AI practices

Responsible use is part of the assignment, and it also matters for a topic like pricing that touches real money.

9.1 Privacy and personal data

We removed the host-related and free-text columns, so the model does not learn from anything that identifies a specific person. The features we kept describe the property and its location, not the host. The dataset is public and was released for a competition, but we still treated it carefully and left those columns out.

9.2 Bias and fairness

Two kinds of bias are worth flagging. The first is geographic. The data is dominated by NYC and LA, so the model has seen far fewer examples from Boston or Chicago and will be less reliable there. The second is that a price model trained on past listings can copy and reinforce existing price gaps between neighbourhoods. If a tool like this were used to recommend prices, it could quietly push already-expensive areas higher and keep cheaper areas low.

9.3 Data leakage

We were careful to avoid leakage. The validation split is separated before any cleaning decision is made. The medians used to fill missing values, the label encoders, and the scaler are all fitted on the training data only and then applied to the validation and test data unchanged.

9.4 Honest reporting

We report MAE, RMSE, and R^2 on a validation set the model was not trained on, not just the best score from training. We are also clear that the official test file has no labels, so we do not invent test metrics for it. We do not claim the model is accurate or finished, and we state plainly that it is off by around 30% and that a third of the price variance is unexplained.

9.5 Real-world risks

The most obvious risk is someone treating the output as a real price today. The data is from 2017 and 2018, and prices have moved a lot since then, so any dollar figure here is out of date. A host who trusted a too-low estimate would lose income, and a guest who trusted a too-high one would overpay. The safe reading of this model is as a rough guide and a learning exercise, not a decision tool.

9.6 Use of AI tools

We used AI assistants for debugging, for explaining error messages, and for improving the grammar of this report. We reviewed and edited the text ourselves, and we can explain every step of the notebook and every decision in this report.

10 Conclusion and future work

10.1 Main findings

We built a full pipeline that predicts Airbnb log prices from property and location features. We loaded the two Kaggle files, explored the training data with statistics, a map, boxplots, and a correlation matrix, cleaned and encoded everything through one shared function, and compared four regression models on a validation split. The Random Forest was the best, reaching an R^2 of about 0.70, with Gradient Boosting close behind and the two linear models tied at about 0.53. The features that matter most are room type, location, and the number of bathrooms, which lines up with common sense about how Airbnb prices work. We then used the Random Forest to generate predictions for the 25,458 unlabeled listings in the test file.

10.2 What we learned

The project taught us how much of data science happens before the model. Cleaning, encoding, and keeping the validation set honest took more thought than fitting the models did. We also saw a concrete case where a feature with almost no linear correlation, the coordinates, turned out to be one of the most useful inputs for a tree model. That gap between linear correlation and tree importance was a good lesson in not trusting a single number. Working with a test file that has no labels also made us think carefully about what evaluation actually means.

10.3 Future work

There is plenty of room to improve. The `amenities` and `description` text could be mined with NLP to add features about what each listing offers. The neighbourhood could be brought back with a smarter encoding such as target encoding instead of being dropped. The tree models deserve proper hyperparameter tuning with cross-validation, and Gradient Boosting in particular might overtake the Random Forest once tuned. Finally, joining the listings to demand or seasonality data would attack the part of the price variance our current features cannot reach.

11 Team contribution

All three members took part in the analysis and the writing. The table below shows the main focus of each person. The work was shared and reviewed together, so these are lead roles rather than strict boundaries.

Table 3: Contribution of each team member.

Member	Main contribution
Mathieu	Data loading of both files, the preprocessing function, missing-value handling, the train/validation split, and the code-quality pass on the notebook.
Serghei	Exploratory data analysis and visualization, the Folium map, the correlation study, feature selection, and the EDA sections of the report.
Samy	Model training and comparison of the four models, validation diagnostics, feature importance, test-file predictions, and the responsible AI and conclusion sections.

Coding was done together in shared sessions, and each member can explain the whole pipeline for the presentation and the Q&A.

Acknowledgements

We would like to thank our instructors, Ajan Aye Hninn Khine and Ajan Naveed Sultan, for their guidance throughout this course. Their lectures gave us the foundations we needed to carry this project from start to finish, and their feedback helped us think more carefully about how we clean data, choose models, and report results honestly.

As exchange students from ESIEA, we also want to say how much we appreciated the welcome we received at KMUTT and the chance to learn data science with Python in such a supportive environment. Thank you for your time and your patience.

References

- [1] Rudy Mizrahi. *AirBnB Listings in Major US Cities (Deloitte ML Competition)*. Kaggle. <https://www.kaggle.com/datasets/rudymizrahi/airbnb-listings-in-major-us-cities-deloitte->
- [2] F. Pedregosa et al. *Scikit-learn: Machine Learning in Python*. Journal of Machine Learning Research, 12:2825–2830, 2011. <https://scikit-learn.org>
- [3] The pandas development team. *pandas documentation*. <https://pandas.pydata.org>
- [4] M. Waskom. *seaborn: statistical data visualization*. <https://seaborn.pydata.org>
- [5] J. D. Hunter. *Matplotlib: A 2D Graphics Environment*. Computing in Science & Engineering, 9(3):90–95, 2007. <https://matplotlib.org>
- [6] Folium contributors. *Folium: Python data, Leaflet.js maps*. <https://python-visualization.github.io/folium>